

ECE 1175

Embedded Systems Design

Lab 5 – CPU Frequency Control over Raspberry Pi

ECE 1175 – Lab 5 (the last one)

- **Prerequisites**

- CPU frequency vs. power consumption
- CPU frequency governor

- **Manipulating CPU Freq. on Raspberry Pi**

- Access CPU frequency info
- Change CPU frequency
- Lab task 1: warm-up of CPU freq. manipulation

- **Customize CPU Freq. Governor**


- “schedutil” governor
- Lab task 2: implement “schedutil” from userspace

Prerequisites

■ CPU Frequency vs. Power Consumption

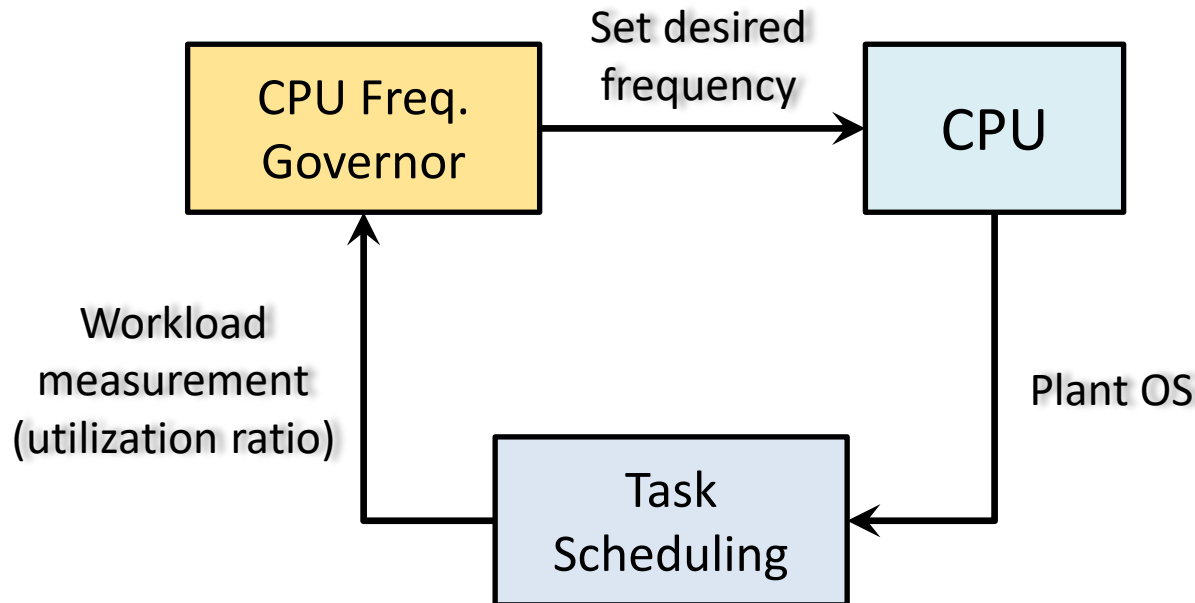
Dynamic CPU power consumption is proportional to the CPU frequency.

$$P_{dyn} = CV^2f$$



Prerequisites

- **CPU Frequency Governor**
 - Control CPU frequency based on different workloads
 - Basic idea: heavy load → high freq. Light load → low freq.



Access CPU Frequency Info

■ Sysfs

- Linux kernel provides an interface via sysfs pseudo filesystem.
- We can get access under:
/sys/devices/system/cpu/cpu[no.]/cpufreq/

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0 $ cd cpufreq
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ ls
affected_cpus          related_cpus           scaling_governor
cpuinfo_cur_freq       scaling_available_frequencies scaling_max_freq
cpuinfo_max_freq       scaling_available_governors scaling_min_freq
cpuinfo_min_freq       scaling_cur_freq       scaling_setspeed
cpuinfo_transition_latency scaling_driver          stats
```

Access CPU Frequency Info

- **Useful CPU Frequency Info**

- `cpuinfo_max_freq` & `cpuinfo_min_freq`
- `cpuinfo_cur_freq`

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat cpuinfo_max_freq
1500000
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat cpuinfo_min_freq
600000
```

Note: the frequency unit is **kHz** by default.

Access CPU Frequency Info

- **Useful CPU Frequency Info**
 - scaling_available_governors
 - scaling_available_frequencies

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat scaling_available_gove  
rnors  
conservative ondemand userspace powersave performance schedutil  
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat scaling_available_freq  
uencies  
600000 750000 1000000 1500000
```

Note: check the scaling available frequencies on Raspberry Pi.

Change CPU Frequency

- We can manually change CPU freq. to any **available** values. To do so,
 - Switch the governor to the **userspace** mode
 - Write your freq. value to **scaling_setspeed** file

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ sudo su
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo userspace > scaling_governor 1. Switch to userspace.
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# cat scaling_governor
userspace
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# cat cpuinfo_cur_freq
600000
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo 1000000 > scaling_setspeed 2. Write any of available freq. values.
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# cat cpuinfo_cur_freq
1000000
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0#
```

Note: you must run these commands as **root**.

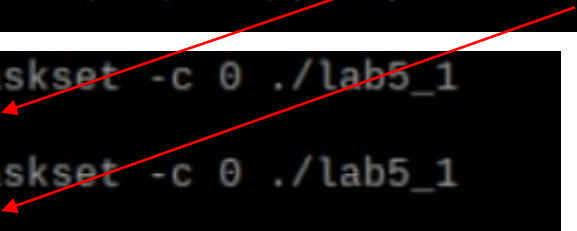
Lab Task 1

■ A Warm-up of CPU Freq. Manipulation

- Get familiar with how to access CPU freq. info
- Measure runtime of the test program under different CPU frequencies.

```
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo userspace > scaling_governor
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo 1500000 > scaling_setspeed
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo 600000 > scaling_setspeed
```

```
pi@raspberrypi:~$ sudo taskset -c 0 ./lab5_1
Execution time: 2.670180s
pi@raspberrypi:~$ sudo taskset -c 0 ./lab5_1
Execution time: 6.813416s
```




Note: Run your program on single core.

Schedutil Governor

■ Schedutil

- Already implemented in the kernel
- For details: <https://lkml.org/lkml/2016/3/17/420>

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat scaling_available_governors  
conservative ondemand userspace powersave performance schedutil
```



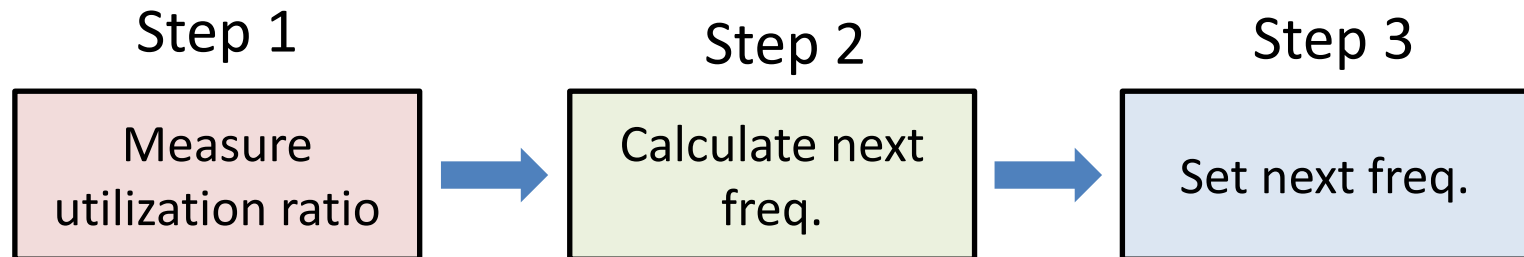
If utilization is frequency-invariant, schedutil says

$$\text{desired_freq} = 1.25 * \text{max_freq.} * \text{util_ratio}$$

Next freq. is the available frequency **closest** to desired_freq

Lab Task 2

- **Write your own program to Implement schedutil from userspace**



- **We are very clear about how to do step 2 and 3.**
 - Step 2: just apply a formula
 - Step 3: write the value to scaling_setspeed file

But how to get utilization ratio in real-time?

Lab Task 2

- **Get per-CPU utilization ratio in real-time**
 - Read /proc/stat file

[illegible]

Lab Task 2

■ Get per-CPU utilization ratio in real-time

- Interpret entries in `/proc/stat`
- For details: [Link](#)

For each `cpu[no.]`,

**Total idle time since boot:
sum 4th & 5th columns**

```
cpu 51869 0 6607 4803123 4126 0 228 0 0 0
cpu0 28376 0 2787 1157054 1996 0 193 0 0 0
cpu1 8942 0 1250 1214468 635 0 22 0 0 0
cpu2 6293 0 1412 1216430 717 0 8 0 0 0
cpu3 8258 0 1158 1215171 776 0 5 0 0 0
```

**Total time since boot: sum 1st to 8th
columns**

The very first "cpu" line aggregates the numbers in all of the other "cpuN" lines. These numbers identify the amount of time the CPU has spent performing different kinds of work. Time units are in USER_HZ (typically hundredths of a second). The meanings of the columns are as follows, from left to right:

- user: normal processes executing in user mode
 - nice: niced processes executing in user mode
 - system: processes executing in kernel mode
 - idle: twiddling thumbs
 - iowait: In a word, iowait stands for waiting for I/O to complete. But there are several problems:
 1. Cpu will not wait for I/O to complete, iowait is the time that a task is waiting for I/O to complete. When cpu goes into idle state for outstanding task io, another task will be scheduled on this CPU.
 2. In a multi-core CPU, the task waiting for I/O to complete is not running on any CPU, so the iowait of each CPU is difficult to calculate.
 3. The value of iowait field in `/proc/stat` will decrease in certain conditions.
- So, the iowait is not reliable by reading from `/proc/stat`.
- irq: servicing interrupts
 - softirq: servicing softirqs
 - steal: involuntary wait
 - guest: running a normal guest
 - guest_nice: running a niced guest

Lab Task 2

- **Get per-CPU utilization ratio in real-time**
 - Calculate utilization ratio for each cpu core

```
// Get accumulative t_total & t_usage
```

$$t_{total} = user + nice + system + idle + iowait + irq + softirq + steal$$
$$t_{idle} = idle + iowait$$
$$t_{usage} = t_{total} - t_{idle}$$

```
// Get real-time util ratio (from moment  $\tau_1$  to  $\tau_2$ )
```

$$\Delta t_{total} = t_{total}(\tau_2) - t_{total}(\tau_1)$$
$$\Delta t_{usage} = t_{usage}(\tau_2) - t_{usage}(\tau_1)$$
$$\%Util = \frac{\Delta t_{usage}}{\Delta t_{total}} \times 100\%$$

One more step: find the **highest** util ratio across all cores as the final util ratio fed to the formula.

Lab Task 2

- **Use programming language of your choice**
 - C/C++: faster, most familiar
 - Python: easier to code
 - Others ...
- **To demonstrate your implementation**
 - Switch to userspace mode
 - Keep running your governor program in the background
 - Run test programs with different workloads
 - TA will check how cpu frequency changes
 - Submit your code to Canvas

Lab Task 2

First run your governor
(use sudo or run in root)

Then run your test program
(e.g., your task1 program)

Check current frequency

Lab Task 2

The image shows three terminal windows on a Raspberry Pi. The top-left window displays CPU statistics: CPU0 at 11.211111%, CPU1 and CPU3 at 0.100000%, CPU2 at 22.322222%, max util at 22.322222%, and next freq at 600000. The top-right window shows the execution of a script in the directory ./lab5_1, with an execution time of 2.766234s. The bottom window, titled pi@raspberrypi: /sys/devices/system/cpu/cpufreq/policy0, shows the user navigating to this directory and listing files. The output of 'ls' lists various scaling parameters. Two red annotations are present: 'busy' with an arrow pointing to the value 1500000 for scaling_cur_freq, and 'idle' with an arrow pointing to the value 600000 for scaling_cur_freq.

```
pi@raspberrypi: ~  
File Edit Tabs Help  
CPU0: 11.211111%  
CPU1: 0.100000%  
CPU2: 22.322222%  
CPU3: 0.100000%  
max util: 22.322222%  
next freq: 600000
```

```
pi@raspberrypi: ~  
File Edit Tabs Help  
pi@raspberrypi:~ $ ./lab5_1  
Execution time: 2.766234s pi@raspberrypi:~ $
```

```
pi@raspberrypi: /sys/devices/system/cpu/cpufreq/policy0  
File Edit Tabs Help  
pi@raspberrypi:~ $ cd /sys/devices/system/cpu/cpufreq/policy0  
pi@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0 $ ls  
affected_cpus      related_cpus      scaling_governor  
cpuinfo_cur_freq   scaling_available_frequencies scaling_max_freq  
cpuinfo_max_freq   scaling_available_governors  scaling_min_freq  
cpuinfo_min_freq   scaling_cur_freq   scaling_setspeed  
cpuinfo_transition_latency scaling_driver      stats  
pi@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0 $ cat scaling_cur_freq  
1500000  
pi@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0 $ cat scaling_cur_freq  
600000  
pi@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0 $
```

Thanks for your participation
in ECE 1175 lab!