

ECE 1175
Embedded Systems Design
Final Review

Wei Gao

Final Exam

- **When**: 4/26 Mon 8:20AM – 9:50AM
 - Students with disabilities will have 50% extra time
- 25% of your final grade
- **What about**: Everything covered in this course
 - Closed-book, closed-notes, no laptop, no discussion
 - All included in class slides (as well as textbook, assigned papers)
 - Distribution: 40% on the first half of semester, 60% on the second half of semester
 - Will include some programming questions
 - More questions than the midterm exam
 - Make your answers short to include only key points
 - Do answer every question, DON'T leave blank

Logistics of Online Exam

- Online proctoring via the lecture zoom link
 - Everyone with your video on throughout the exam
 - Make sure that we can see not only your face, but also your hands
- Exam procedure
 - The exam sheet will be distributed online via zoom right before the exam starts
 - You will need to write your answers on paper and scan your answers back
 - Recommended software: Adobe scan via smartphone
 - Make sure that you settle a scanning method and get it ready before the exam!
 - After the exam
 - Stay online with video on
 - You will need to scan your answers and send them out right after the exam
 - You can only log off after your answers are received
 - You have to make sure your answers are received within 30 minutes after the exam

First half of the semester

- Introduction to embedded systems
 - Real-time, low power, small memory footprint, low cost
 - Soft/hard real-time system
- Design methodology
 - Microprocessors, FPGA and ASIC
 - Design procedure and example: GPS
- Microprocessors
 - von Neumann vs. Harvard, RISC vs. CISC, SHARC, ARM7
- CPUs, I/O, interrupt
 - Busy-wait I/O, interrupt-based I/O, interrupt mechanism
- Caches and memory
 - Memory system, average memory access time in multi-level cache
 - Cache organization: direct-mapped, N-way set-associative
- Real-time scheduling

Definition

- **Embedded system**: any device that includes a computer but is not itself a general-purpose computer.
- System characteristics
 - Non-functional requirements: Real-time, Low power, Small memory footprint, Low cost
 - Hard vs. soft real-time

Alternative Technology

- Application-Specific Integrated Circuits (ASICs)
- Microprocessors
- Field-Programmable Gate Arrays (FPGAs)
- Why should we use microprocessors?
 - Reprogrammability and low development cost >> low performance/watt

Microprocessors

- Performance

- Con: Programmable architecture is fundamentally slow!
 - Fetch, decode instructions
- Pro: Highly optimized architecture and manufacturing
 - Pipelines; cache; clock frequency; circuit density; manufacturing technology


- Power

- Processors perform poorly in terms of performance/watt!
- Power management can alleviate the power problem.

- Flexibility, development cost and time

- Let software do the work!

Design Methodologies



Microprocessors

- von Neumann
 - Same memory holds data, instructions.
 - A single set of address/data buses between CPU and memory
- Harvard
 - **Separate** memories for data and instructions.
 - Two sets of address/data buses between CPU and memory
- RISC vs. CISC
 - CISC: Many addressing modes and instructions; High code density.
 - RISC: Compact, uniform instructions: facilitate pipelining, poor memory footprint

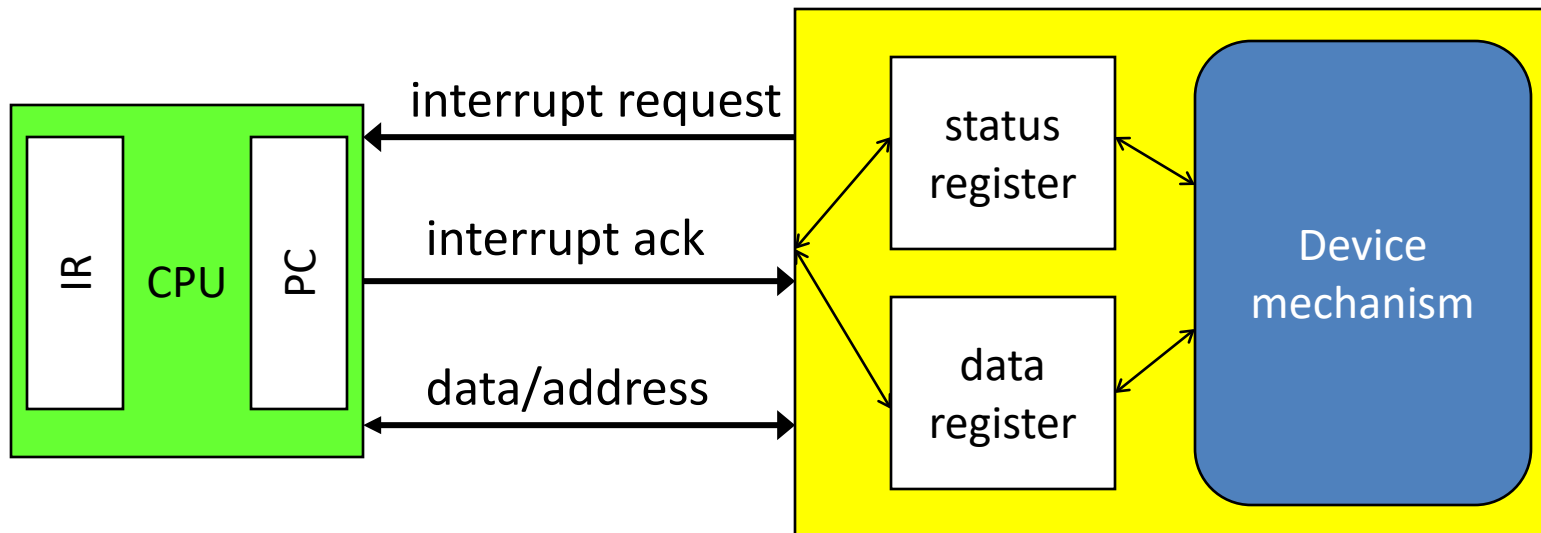
Busy-Wait I/O Programming

- Simplest way to program I/O devices.
 - Devices are usually slower than CPU and require more cycles
 - CPU has to wait for device to finish before starting next one
 - Use `peek` instruction to test when device is finished
 - Test-and-set

```
//send a string to device using Busy-Wait handshaking
current_char = mystring;
while (*current_char != '\0') {
    //send character to device (data register)
    poke(OUT_CHAR, *current_char);
    //wait for device to finish by checking its status
    while (peek(OUT_STATUS) != 0);
    //advance character pointer to next one
    current_char++;
}
```


Interrupt-based I/O

- Busy-wait is very inefficient.
 - CPU can't do other work while testing device.
 - Hard to do simultaneous I/O.
- Interrupts allow to change the flow of control in the CPU.
 - Call interrupt handler (i.e. device driver) to handle device.




Microprocessor Bus

- Bus is a set of wires and a protocol for the CPU to communicate with memory and devices
- Five major components to support reads and writes



Typical Bus Access

- Timing diagram syntax:
 - Tri-state: Constant value (0/1), stable, changing, unknown.



I/O Interfaces

- Parallel v.s. Serial
 - Parallel
 - Wider bandwidth
 - More wires indicate more overhead
 - Simple I/O operation
 - Serial
 - 1-bit transfer per time unit
 - Less wires indicate less overhead
 - Complex I/O protocol
- Serial I/O interfaces
 - SPI
 - I2C
 - Master/slave mode



ARM Architecture

- Dominant architecture for embedded systems



ARM Instruction Sets

- ARM instructions can be made to execute conditionally by postfixing them with the appropriate condition code field.
 - This improves code density *and* performance by reducing the number of forward branch instructions.

```
CMP    r3,#0
BEQ    skip
ADD    r0,r1,r2
skip
```



```
CMP    r3,#0
ADDNE  r0,r1,r2
```


Memory System and Caches

- Memory is slower than CPU
 - CPU clock rates increase faster than memory
- Caches are used to speed up memory
 - Cache is a small but fast memory that holds copies of the contents of main memory
 - More expensive than main memory, but faster
- Memory Management Units (MMU)
 - Memory size is not large enough for all application?
 - Provide a larger virtual memory than physical memory

Memory Devices

- Types of memory devices
 - RAM (Random-Access Memory)
 - Address can be read in any order, unlike magnetic disk/tape
 - Usually used for data storage
 - DRAM vs. SRAM.
 - ROM (Read-Only Memory)
 - Usually used for program storage
 - Mask-programmed vs. field-programmable.

Key OS Issues

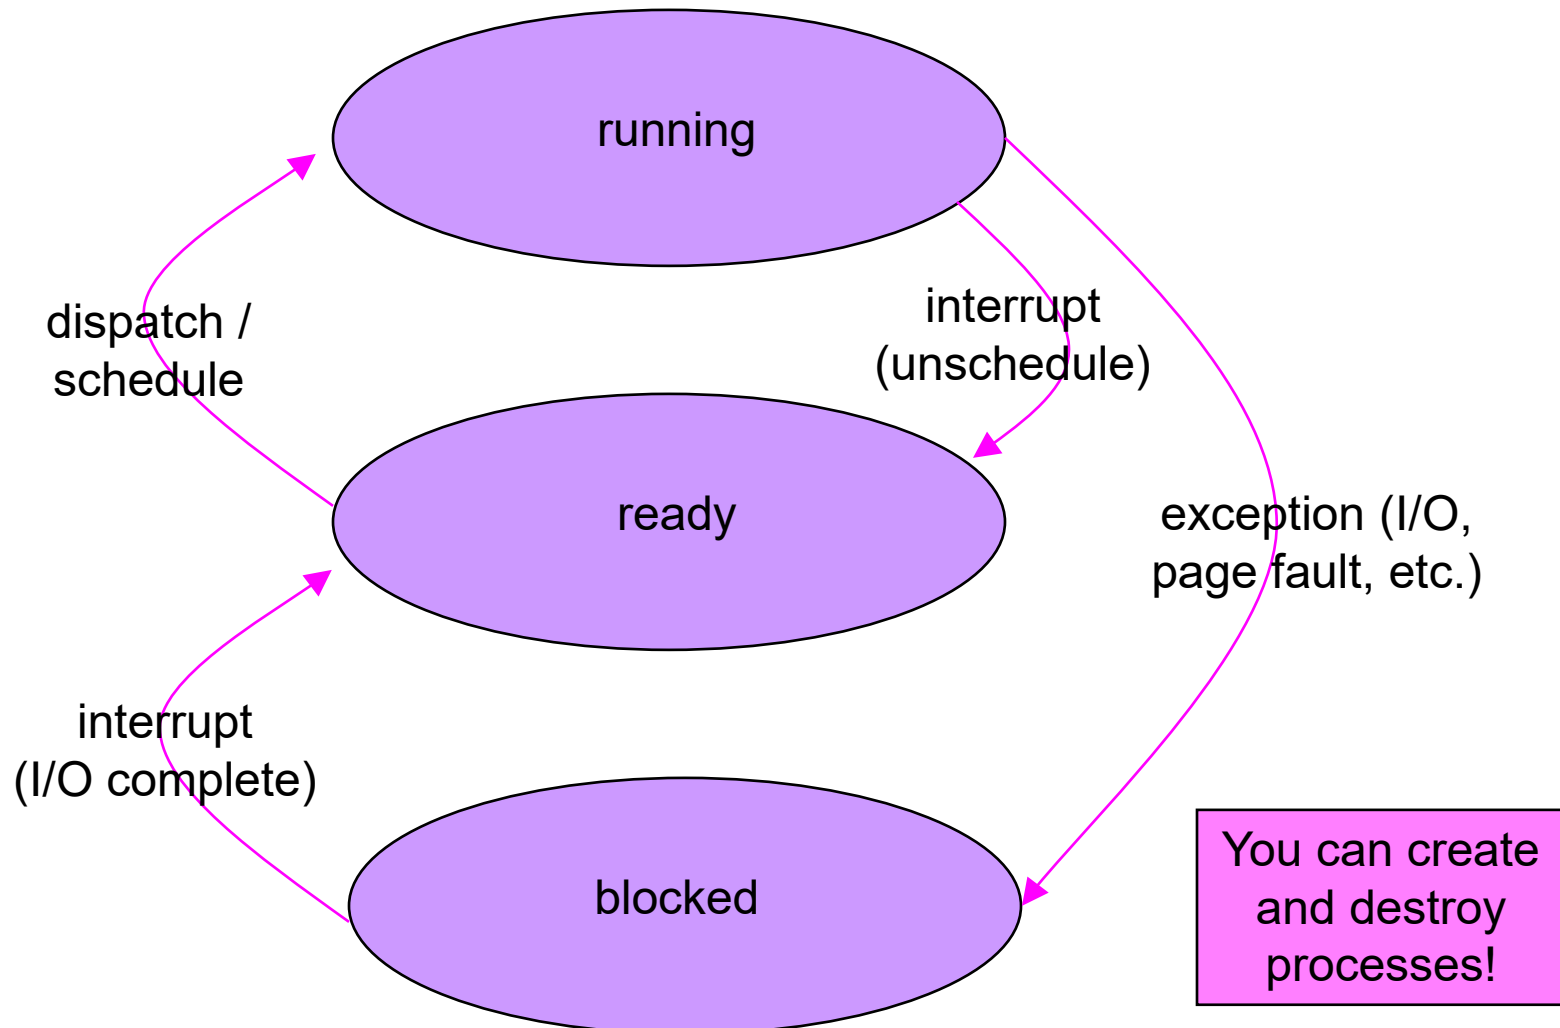
■ Multi-Programming

- keeps multiple runnable jobs loaded in memory at once
- overlaps I/O of a job with computing of another
 - while one job waits for I/O completion, OS runs instructions from another job
- goal: **optimize system throughput**
 - perhaps at the cost of response time...

■ Timesharing


- multiple terminals into one machine
- each user has illusion of entire machine to him/herself
- **optimize response time**, perhaps at the cost of throughput
- Timeslicing
 - divide CPU equally among the users

A Process's Lifecycle



Deadlock

- Resource graph
 - A deadlock exists if there is an *irreducible cycle* in the resource graph



- R1 is held by
- - → is waiting for R1
- R2 is held by
- - → is waiting for R2

Second Half of the semester

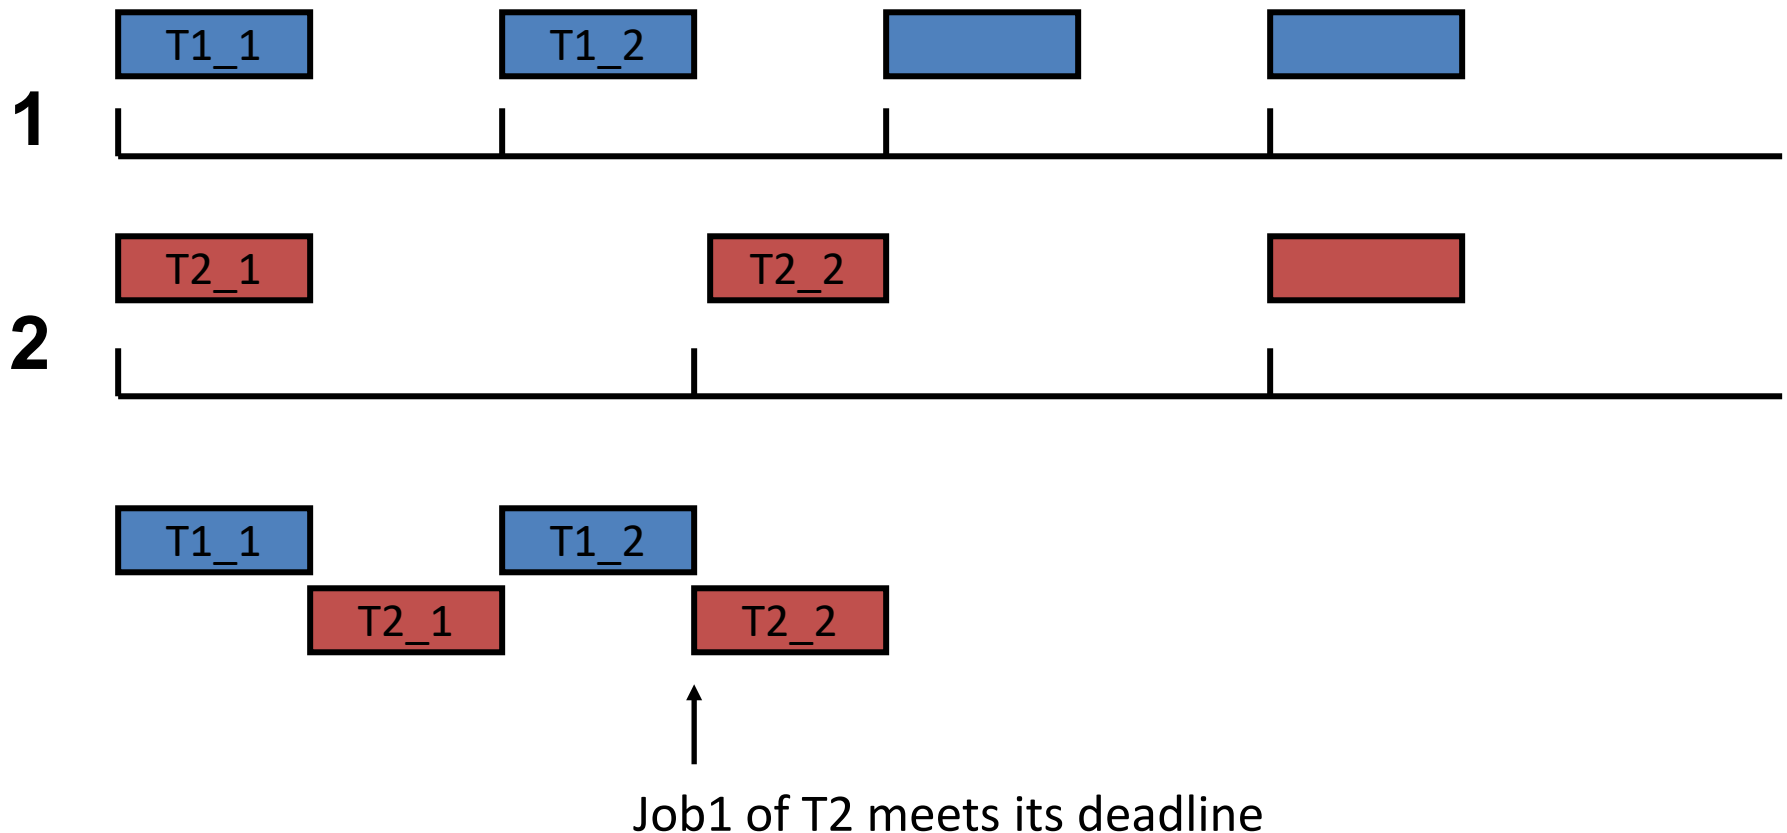
- Real-Time Scheduling
- Program optimizations
- Power management
- Safety and Reliability

Real-time Scheduling

- Terminologies and timing parameters
 - Task, job, subtask
- Metrics to evaluate scheduling algorithms
 - Schedulability, overhead
- Optimal scheduling algorithms
 - When relative deadline = period: RMS, EDF, utilization bound
 - When relative deadline < period: EDF, processor demand analysis
- CPU utilization analysis and bound
- Priority inversion
 - Sources, unbounded priority inversion, priority inheritance
- End-to-end scheduling framework
 - Task allocation: bin packing
 - Synchronization protocol: greedy protocol, release guard
 - Subdeadline assignment: ultimate deadline, proportional deadline


RMS Meeting the Deadline

- $T1 = (10, 20)$, $T2 = (10, 30)$, utilization is 83%




EDF Meeting a Deadline

- $T1 = (10,20)$, $T2 = (15,30)$, utilization is 100%




Priority Inversion

critical section




Unbounded Priority Inversion



Solution: Priority Inheritance

- Let the low-priority task inherit the priority of the blocked high-priority task.

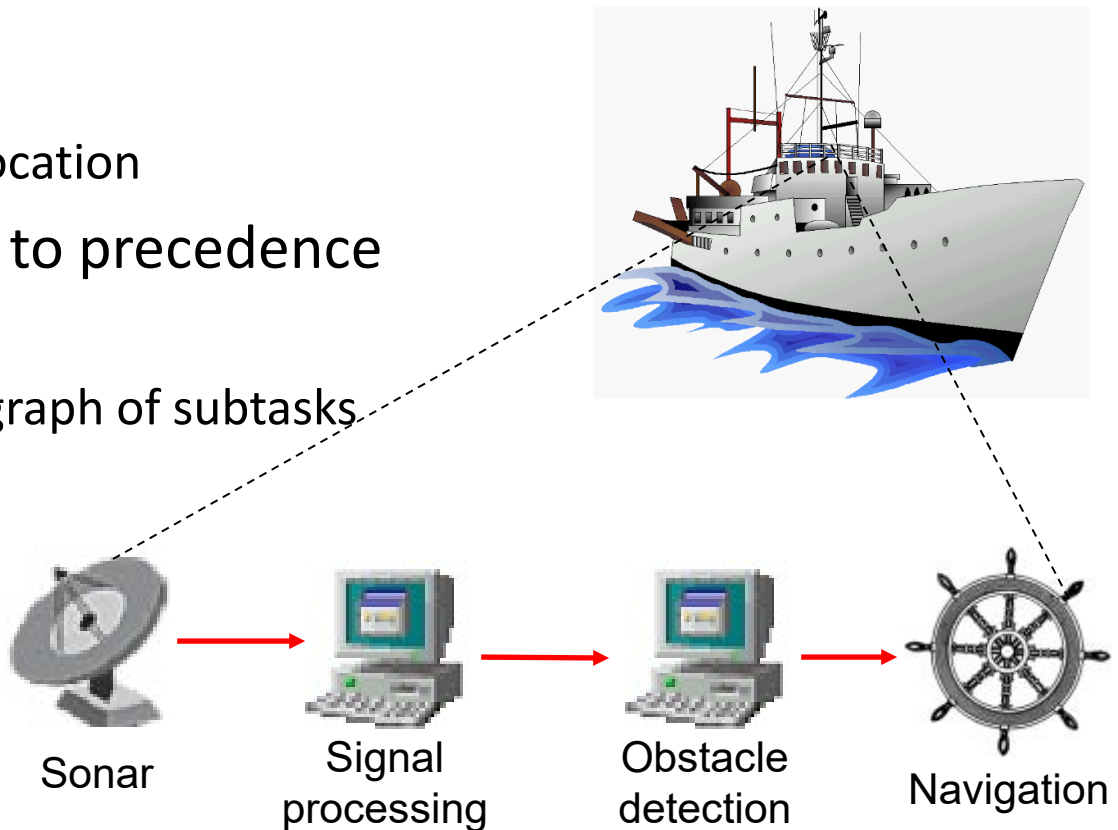


Multi-Processor Systems

- Tight coupling among processors.
- Communicate through shared memory and on-board bus.
- Scheduled by a common scheduler/OS.
 - Global scheduling
 - Partitioned scheduling
- States of all processors available to each other.

End-to-End Task Model

- An (end-to-end) task is composed of multiple subtasks running on multiple processors
 - Message/event
 - Remote method invocation
- Subtasks are subject to precedence constraints
 - Task = a **chain**/tree/graph of subtasks
 - E.g. ship navigation




End-to-End Scheduling Framework

1. Task allocation
2. Synchronization protocol
3. Subdeadline assignment
4. Schedulability analysis



Greedy Protocol

- After a subtask is finished, the next subtask starts **immediately**
- Release job $J_{i,j;k}$ **as soon as** $J_{i,j-1;k}$ is completed
- Subsequent subtasks may **not** be periodic under a greedy protocol
 - Difficult for schedulability analysis
 - High-priority tasks arrive early \rightarrow high worst-case response time for lower-priority tasks



Greedy Protocol Illustrated

- Task: (C,P)




Release Guard

- After a subtask is finished, the next subtask may wait for a while before release
- Every subtask (if not a first subtask) has a release guard, which
 - waits for the preceding subtask for a result/event
 - then releases the job
 - at the point of exact one period from the last release time (Rule1)
OR
 - whenever the processor becomes idle (Rule 2)
- Release guard strategy improves worst response time without affecting schedulability

Release Guard Illustrated

■ Task: (C,P)



Basic Compilation Optimization

- Expression simplification
- Dead code elimination
- Function inlining
- Loop optimizations
- Register allocation

1. Expression Simplification


- Algebraic:
 - $a*b + a*c \Rightarrow a*(b+c)$
 - Three operations change to two operations
- Constant folding:
 - `for (i=0; i<8+1; i++)`
 \Rightarrow `for (i=0; i<9; i++)`
- Strength reduction of constant multiplication
 - $a*2 \Rightarrow a \ll 1$; $a*3 \Rightarrow a*(2+1) \Rightarrow a \ll 1 + a$;
 - $a*1000 \Rightarrow a*(1024-16-8) \Rightarrow a \ll 10 - a \ll 4 - a \ll 3$
 - SHIFT, ADD, SUB take one cycle but multiplication takes about 10 cycles

2. Dead Code Elimination

- Dead code:

```
#define DEBUG 0  
if (DEBUG) dbg(p1);
```

- Can be eliminated by analysis of control flow, constant folding



Function inlining

```
int foo(a,b,c) { return a + b - c; }  
z = foo(w,x,y);
```



```
z = w + x - y;
```

- An inline function's body is inserted directly (like a **substitution**) in the compiled code at the point where the function is called.
- Improve performance by reducing function call overhead
- “inline” in different cases
 - TinyOS does whole-program inlining

Loop Optimizations

- Loops are good targets for optimization.
- Basic loop optimizations:
 - Code motion;
 - Reduce loop overhead: loop unrolling
 - Increase opportunities for pipelining and parallelism: loop fusion

Register Allocation

- Processor registers
 - A very small amount of very fast computer memory
 - Used to speed the execution of computer programs
 - Provides quick access to **most commonly** used values
 - **Memory hierarchy**: register – cache – main memory – disk
- Reduce the number of used registers
 - Fit more frequently used variables in registers
 - Load once, use many times

Register Lifetime Graph



no. of needed registers = 5

1. $w = a + b;$

2. $x = c + w;$

3. $y = c + d;$

4. $z = a - b;$



means this variable should be loaded to register

After Rescheduling

no. of needed registers = 4

1. **w** = **a** + **b**;

2. **z** = **a** - **b**;

3. **x** = **c** + **w**;

4. **y** = **c** + **d**;

a				
b				
c				
d				
w				
x				
y				
z				
	1	2	3	4

Cannot change dependencies between instructions!

Power Management

- Hardware support
 - CMOS features: voltage drops, toggling, leakage
 - Clock gating, supply shutdown, dynamic voltage scaling
- Power management policy
 - Dynamic power management
 - Power state machine, break-even time T_{BE}
 - Energy saving calculation based on a known idle time
 - Predictive techniques
 - Metrics of prediction quality: safety and efficiency
 - Fixed timeout vs. predictive shutdown/wakeup
- Power manager
 - Advanced Configuration and Power Interface (ACPI)
- Holistic approach
 - Memory system, cache behavior

CMOS Power Consumption

- All digital systems are built with CMOS
 - CMOS: Complementary Metal-Oxide Semiconductor
 - A common technology for constructing integrated circuits
- Voltage drops
 - Power consumption of a CMOS is proportional to the square of the power supply voltage (V^2).
- Toggling
 - CMOS uses higher power when having more activity
- Leakage
 - Even when CMOS is inactive, some charge leaks out of the circuit's nodes

General Power-Saving Features

To deal with toggling

- ➡ Run at lower clock frequency (slower)
- ➡ Reduce activity by disabling function units when not in use.

To deal with leakage

- ➡ Eliminate leakage current by disconnecting parts from power supply when not in use.

To deal with voltage drops

- ➡ Reduce power supply voltage to the **lowest** level that provides required performance
 - ➡ Pentium MMX: 2.8v
 - ➡ i7: ~1.5v
 - ➡ Cortex A8: 1.2v


Dynamic Voltage Scaling

■ Why **voltage scaling**?

- Power $\propto V^2$
 - Reduce power supply voltage \rightarrow save energy
- Lower clock frequency allows lower voltage
 - Tradeoff between performance and battery lifetime
 - Overclocking? Raise voltage for better stability

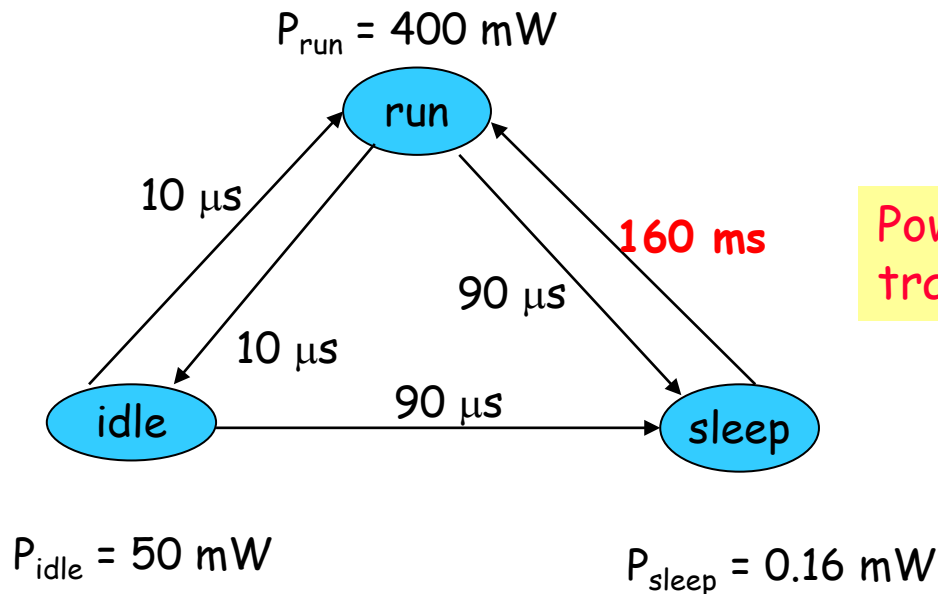
■ Why **dynamic**?

- Power consumption is not a constant.
 - Depending on current workloads.
 - Peak computing rate is usually much higher than average.
- Have to be dynamic to respond to power consumption variations.



Break-Even Time T_{BE}


- T_{BE} of an inactive state is the total time for entering and leaving the state
 - Assumption: transition doesn't cause extra power consumption
- $T_{BE} = T_{TR} = T_{On,Off} + T_{Off,On}$
 - Ex. $T_{BE} = 160 \text{ ms} + 90 \text{ }\mu\text{s}$ for SLEEP in SA-1100



Power consumption during transition $\approx P_{run}$

Break-Even Time When $P_{TR} > P_{On}$

- $T_{BE} = T_{TR} + T_{TR}(P_{TR} - P_{On})/(P_{On} - P_{Off})$
 - $T_{TR}(P_{TR} - P_{On})$: extra energy consumed for transition
 - $/(P_{on} - P_{off})$: the idle time needed to compensate this energy consumption back
- It is easier to save power with a shorter T_{BE}
 - Shorter T_{TR}
 - IDLE in SA-1100
 - Higher difference between $P_{On} - P_{Off}$
 - SLEEP in SA-1100
 - Lower P_{TR}



Energy Saving Calculation

- Given an idle period $T_{\text{idle}} > T_{\text{BE}}$
 - $E_S(T_{\text{idle}}) = (T_{\text{idle}} - T_{\text{TR}})(P_{\text{On}} - P_{\text{OFF}}) + T_{\text{TR}}(P_{\text{On}} - P_{\text{TR}})$
 - $P_{\text{On}} > P_{\text{TR}}$: total = idle saving + transition saving
 - $P_{\text{On}} < P_{\text{TR}}$: total = idle saving - transition cost
- Achievable power saving depends on workload!
 - Distribution of idle periods

Predictive Techniques

- Don't know how long the idle time will be?
- Predict whether idle time $T_{\text{idle}} > T_{\text{BE}}$ based on history
 - Ex: someone takes break once an hour?
 - Predicted event: $p = \{T_{\text{idle}} > T_{\text{BE}}\}$
 - Observed event: o
 - Triggers state transition: RUN -> IDLE/SLEEP

Metrics of Prediction Quality

- **Safety**: conditional probability $\text{Prob}(p | o)$
 - If an observed event happens, what's the probability of $T_{\text{idle}} > T_{\text{BE}}$?
 - Ideally, safety = 1.
- **Efficiency**: $\text{Prob}(o | p)$
 - If $T_{\text{idle}} > T_{\text{BE}}$, what's the probability of successfully predicting it in advance (e.g., o happens)?
- **Overprediction**
 - State transition too much
 - High performance penalty \rightarrow poor safety
- **Underprediction**
 - State transition not enough
 - Wastes energy \rightarrow poor efficiency

Fixed Timeout Policy

- Enter inactive state when the system has been idle for T_{TO} sec.
 - o: $T_{idle} > T_{TO}$
- Wake up in response to activity
- Assumption: If a system has been idle for T_{TO} sec, it is likely that it will continue to be idle for $T_{idle} - T_{TO} > T_{BE}$.

Selecting Best Timeout T_{TO}

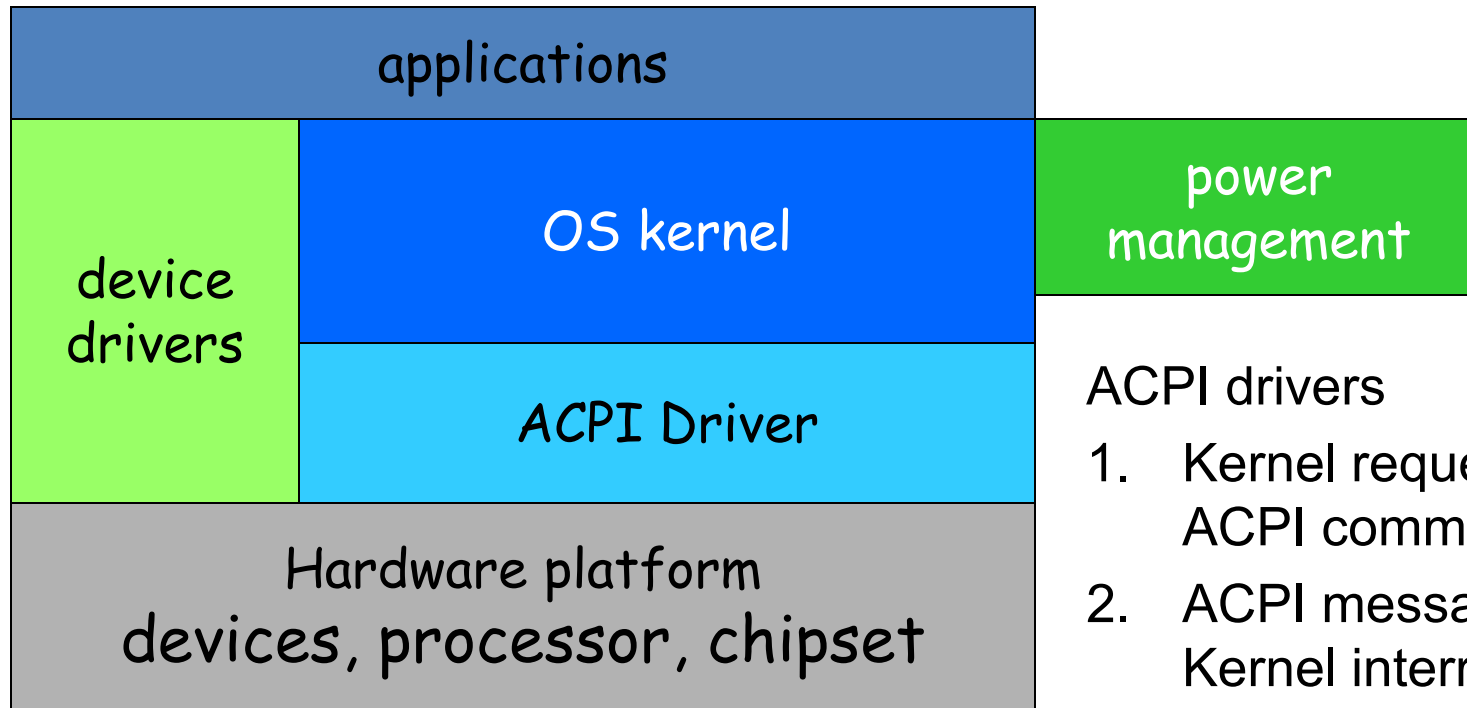
- Increasing T_{TO} improves safety, but reduces efficiency at the same time
 - Safety: $\Pr(T_{idle} > T_{TO} + T_{BE} \mid T_{idle} > T_{TO})$
- Highly workload dependent
- Karlin's result: $T_{TO} = T_{BE}$
 - Energy consumption is at worst twice the energy consumed by an ideal policy

Possible Improvement

- **Predictive shutdown**: shutdown *immediately* when an idle period starts
 - Avoid wasting energy before reaching the timeout threshold
 - More efficient, less safe
- **Predictive wakeup**: wake up before the *predicted* idle time expires, even if no new activity has occurred.
 - Avoid performance penalty for wake up
 - Less efficient, safer


Advanced Configuration and Power Interface (ACPI)

- Open standard for power management services.
 - Proposed by Intel, Microsoft and Toshiba



What is a “Safe” Embedded System?

- Again, the ABS system’s case



- Is it safe?
- Add watch dog between brake and bus
- What does “safe” mean?
- Add mechanical linkage from brake pedal directly to brake
- How can we make it safe?
- Add a third mechanical linkage...

Component vs. System

- Reliability is a component issue
- **Safety** and **Availability** are system issues
- A system can be safe even if it is unreliable!

- If a system has lots of redundancy the likelihood of a component failure (a fault) increases, but so may increase the safety and availability of that system

- Safety and Availability are different and sometimes at odds. Safety may require the shutdown of a system that may still be able to perform its function.
 - A backup system that can fully operate a nuclear power plant might always shut it down in the event of failure of the primary system.
 - The plant could remain available, but it is unsafe to continue operation

Terms

- Safety: assuming acceptable risk
- Hazard: dangerous system state that could cause damage
- Fault: conditions that lead to hazards
- Reliability
 - System is functioning if all components are functioning
 - $P_s(t) = \prod_i P_i(t)$
 - System is functioning
 - $P_s(t) = 1 - \prod_i F_i(t)$
 - $F_i(t) = 1 - P_i(t)$: probability of component failure

Safety Architectures

- Self checking (single channel protected design)
 - Perform periodic checksums on code and data
 - How long does it take?
 - No protection against systematic faults
- Redundancy
 - Homogeneous redundancy
 - Heterogeneous redundancy
- Diversity or Heterogeneity