

ECE 1175
Embedded Systems Design
Real-Time Scheduling - I

Wei Gao

The Scheduling Problem


- Can we meet all deadlines?
 - Must be able to meet deadlines in all cases.
- How much CPU horsepower do we need to meet our deadlines?
- Timing violations: What happens if a process doesn't finish by its deadline?
 - **Hard deadline**: system fails if missed.
 - **Soft deadline**: user may notice, but system doesn't necessarily fail.

Process Scheduling: Embedded vs. General-Purpose

- General-purpose systems
 - e.g., PCs, database servers
 - **Fairness** to all tasks (no starvation)
 - Optimize **throughput**
 - Optimize **average** performance
- Embedded systems
 - Meet all **deadlines**.
 - Fairness or throughput is **not** important
 - Hard real-time: worry about **worst case** performance


Terminologies Used in Scheduling

- **Task**
 - May correspond to a process or thread
 - May be released multiple times
- **Periodic task**
 - Ideal: inter-arrival time = **period**
 - General: inter-arrival time \geq **period**
- **Non-periodic task**
 - Inter-arrival time does not have a lower bound
- **Job**: an instance of a task




Timing Parameters - Task

- Task T_i
 - Period P_i
 - Worst-case execution time C_i
 - **Relative** deadline D_i
 - Usually equal to period



Timing Parameters - Job

- Job J_{ik} (denoted as Ti_k)
 - Release time: time when a job is ready
 - Response time R_i = finish time – release time
 - **Absolute** deadline = release time + D_i
- A job misses its deadline if
 - Response time $R_i > D_i$
 - Finish time > absolute deadline



Metrics to Evaluate Scheduling Algorithms

- **Schedulability**
 - A task set is **schedulable** under a scheduling algorithm if all jobs can meet their deadlines
- **Overhead**
 - Time required for scheduling decision and context switches.

Optimality

A scheduling algorithm S is **optimal** if


- a task set is not schedulable under $S \rightarrow$ it is not schedulable under any other algorithms

CPU Utilization Analysis

- **Utilization** of a processor:


$$U = \sum_{i=1}^n \frac{C_i}{P_i}$$

n: number of tasks on the processor



CPU utilization of T1 = $3/5 = 60\%$

CPU utilization
of a processor



Schedulable Utilization Bound

- Utilization bound U_b
 - All tasks are guaranteed to be schedulable if $U \leq U_b$
 - U is the requested utilization of a task set
- Conditions for scheduling
 - No scheduling algorithm can schedule a task set if $U > 1$
 - $U_b \leq 1$
 - An algorithm is optimal if its $U_b = 1$

Optimal Scheduling Algorithms

- Rate Monotonic Scheduling (RMS)
 - Higher rate ($=1/\text{period}$) \rightarrow Higher priority
 - Optimal preemptive **static** priority scheduling algorithm
- Earliest Deadline First (EDF)
 - Earlier **absolute** deadline \rightarrow Higher priority
 - Optimal preemptive **dynamic** priority scheduling algorithm

Assumptions

- Single processor.
 - All tasks are periodic.
 - Zero context switch time.
 - Relative deadline = period.
 - No priority inversion.
-
- RMS and EDF have been extended to cases with relaxed assumptions

RMS - Rate Monotonic Scheduling

- Common way to assign priorities
- Result from Liu & Layland, 1973 (JACM)
- Simple to understand and implement:

Processes with shorter period are
given higher priority

- E.g.,

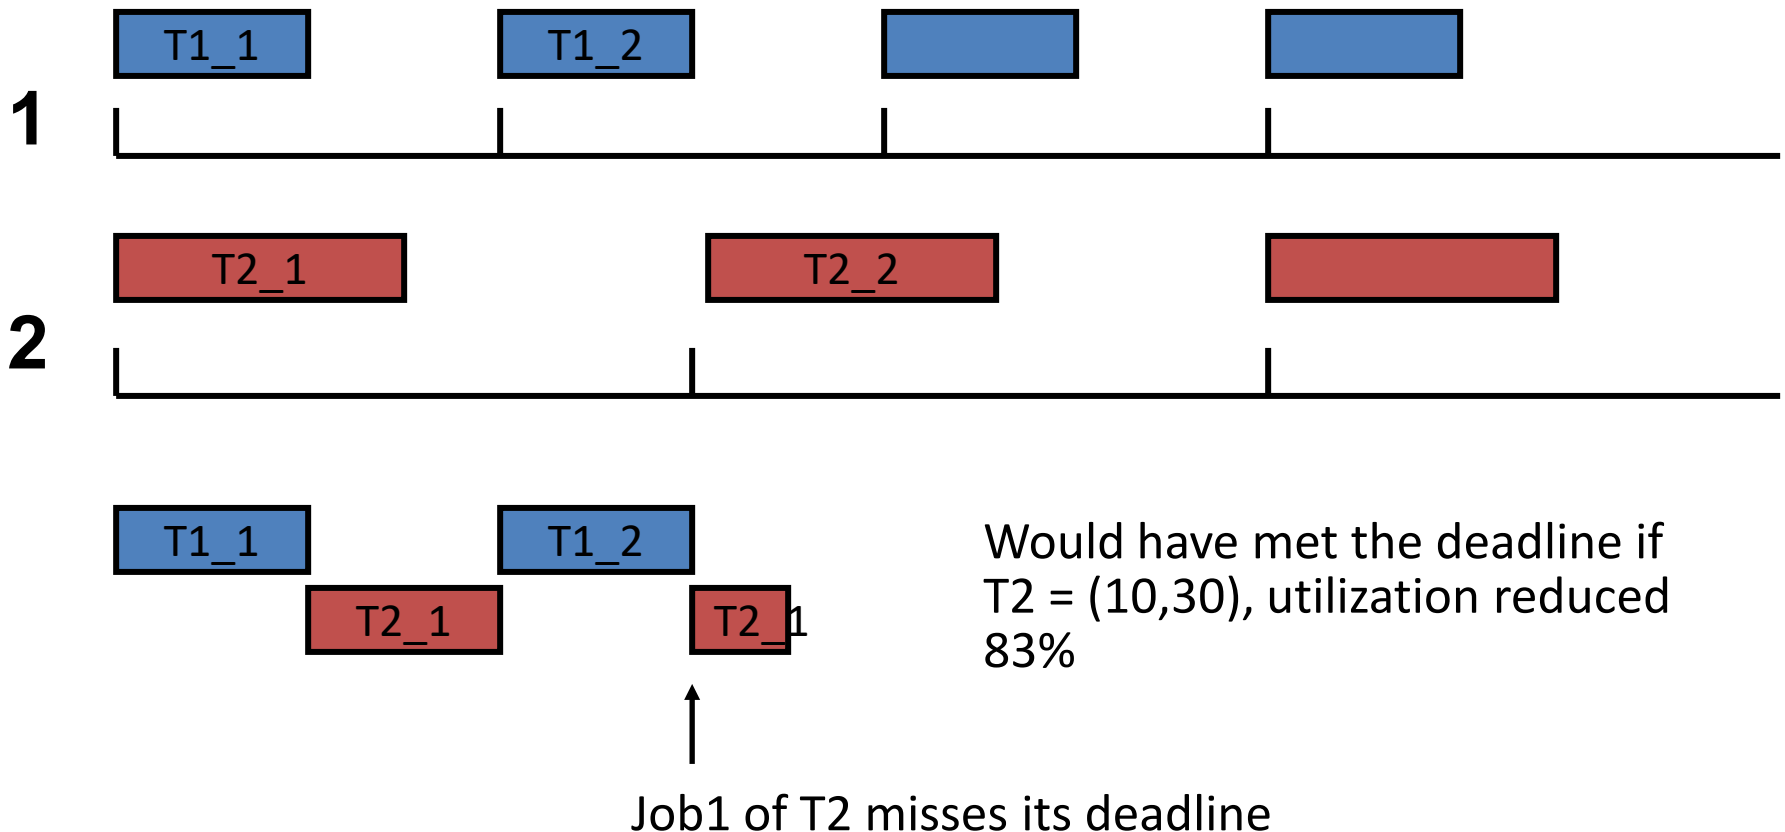
<u>Period</u>	<u>Priority</u>	
10	1	(highest)
12	2	
15	3	
20	4	(lowest)

RMS Utilization Bound

- $U_b(n) = n(2^{1/n} - 1)$
 - n : number of tasks
 - $U_b(1) = 1$
 - $U_b(2) = 0.828$
 - $U_b(n) \geq U_b(\infty) = \ln 2 = 0.693$
- $U \leq U_b(n)$ is a **sufficient** condition, but **not necessary** in general cases.
- $U_b = 1$ if all process periods are **harmonic**, i.e., periods are multiples of each other
 - e.g., 1,10,100

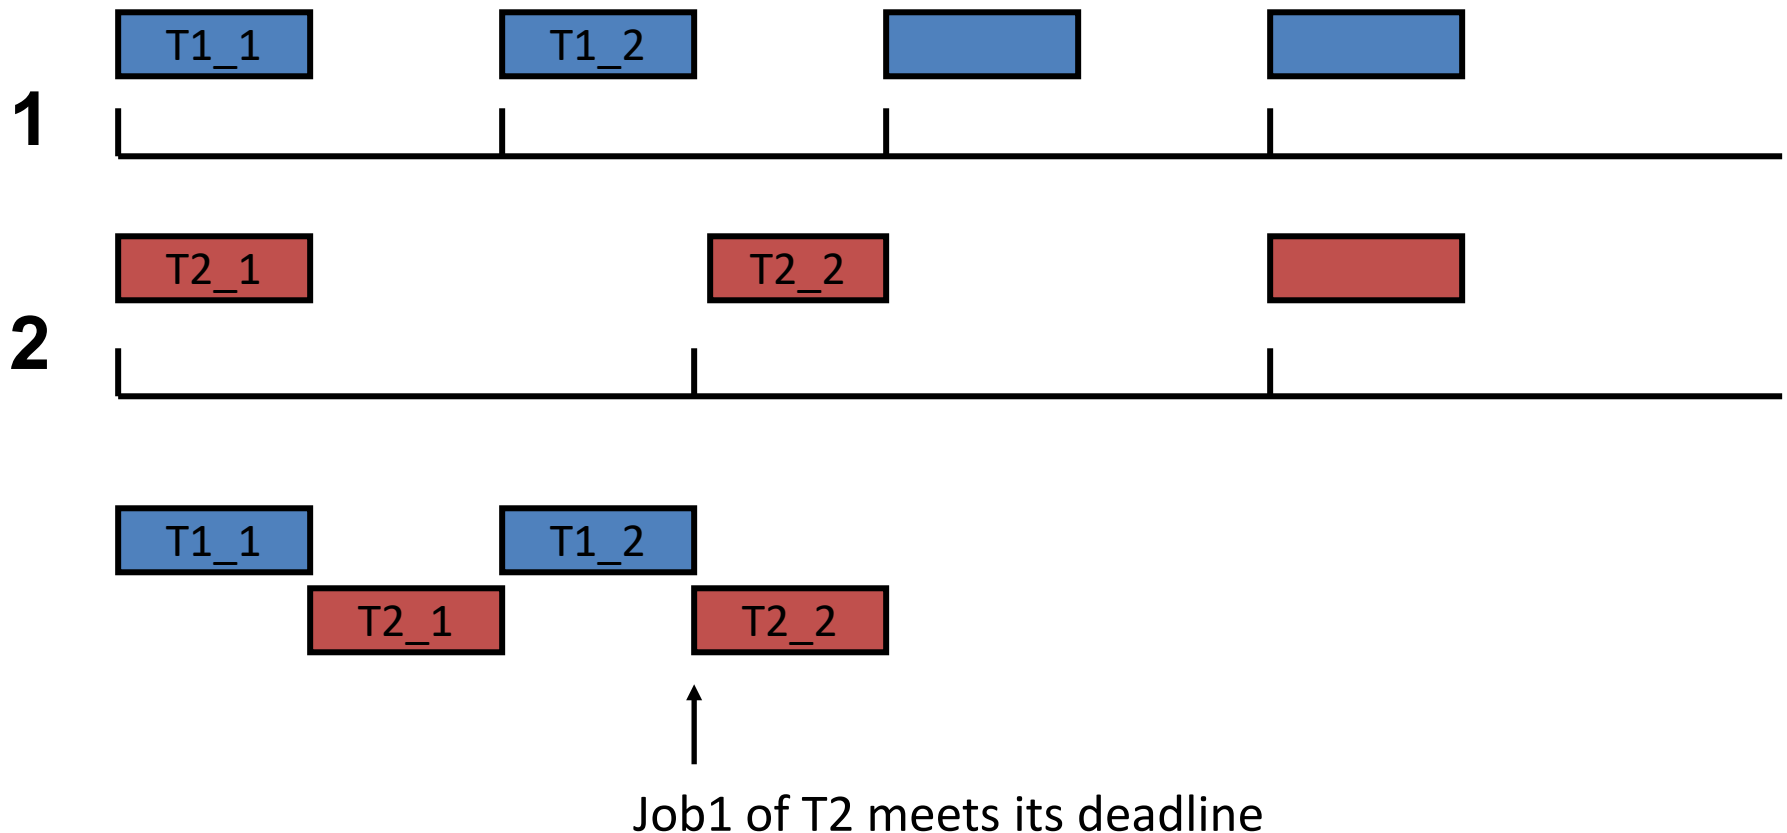
RMS Missing the Deadline

- $T1 = (10, 20)$, $T2 = (15, 30)$, utilization is $100\% > \text{RMS bound}$



RMS Meeting the Deadline

- $T1 = (10, 20)$, $T2 = (10, 30)$, utilization is 83%

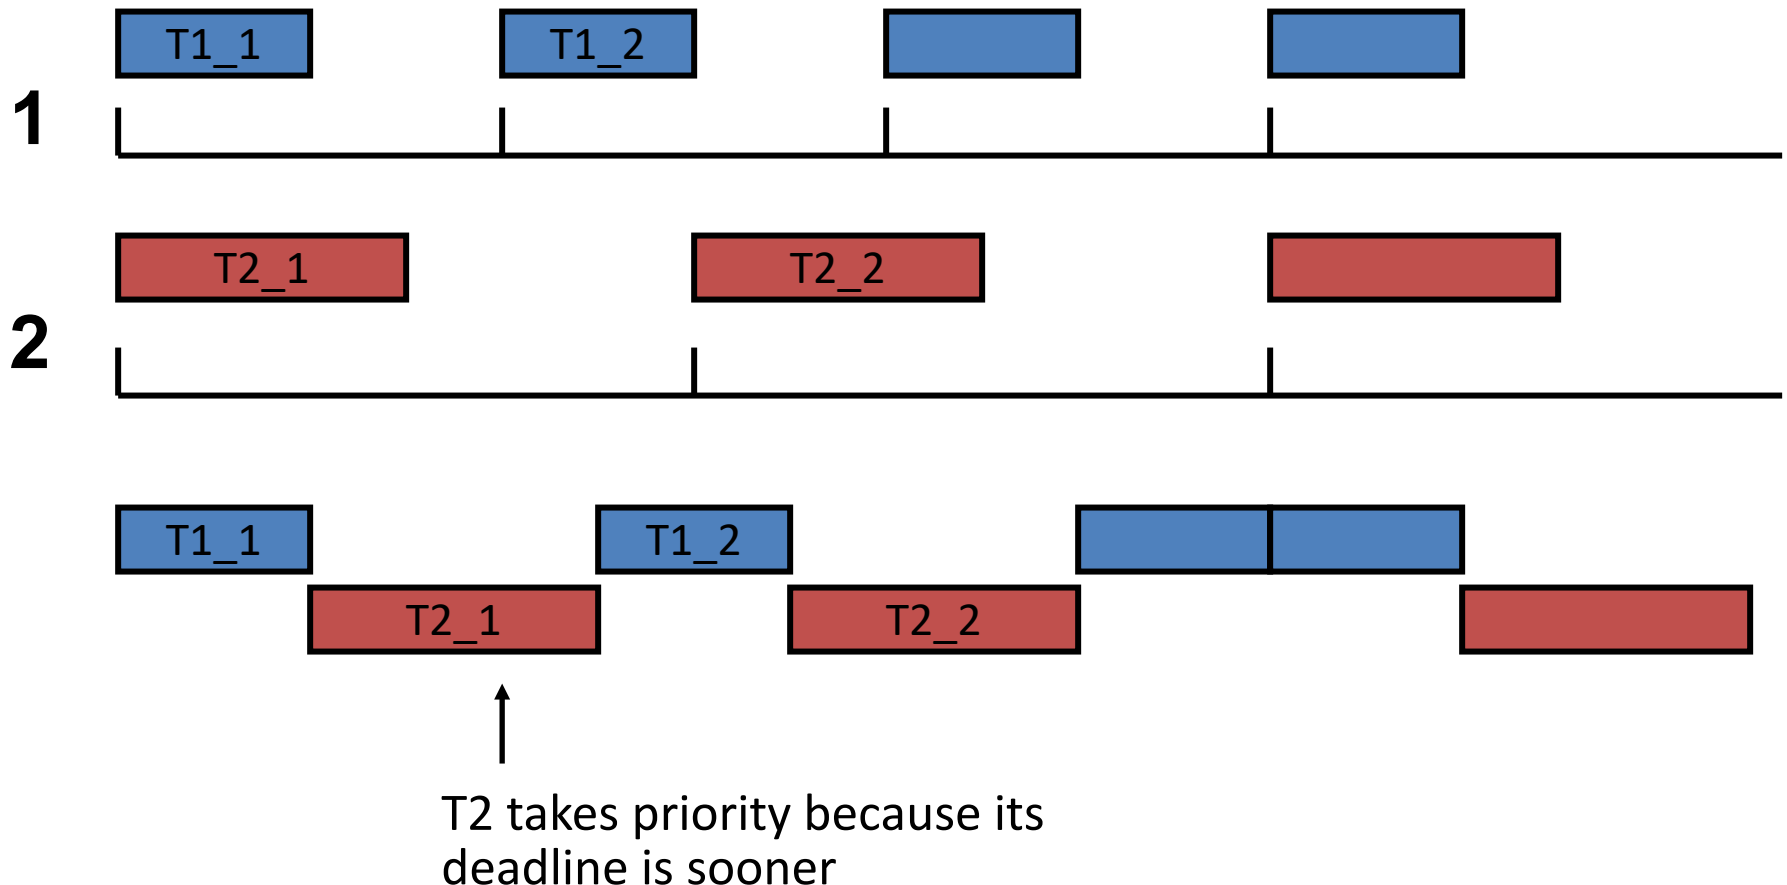


RMS Evaluation

- Schedulability
 - RMS may not guarantee schedulability even when CPU is not fully utilized
- Low overhead
 - When tasks are fixed, priorities are never changed
- Optimal **static** priority scheduling algorithm

EDF Meeting a Deadline

- $T1 = (10, 20)$, $T2 = (15, 30)$, utilization is 100%



EDF Utilization Bound

- $U_b = 1$
- $U \leq 1$ is a **sufficient** and **necessary** condition for schedulability.
- Generally considered **too expensive to use in practice**.
 - Absolute deadline for each job needs to be computed
 - Change process priorities on the fly

EDF Evaluation

- Schedulability
 - EDF can guarantee schedulability as long as CPU is not fully utilized
- Higher overhead than RMS
 - Task priorities may need to be changed online
- Optimal **dynamic** priority scheduling algorithm

Summary

- Terminologies and timing parameters
 - Task, job
- Metrics to evaluate scheduling algorithms
 - Schedulability, overhead
- Optimal scheduling algorithm: Rate Monotonic Scheduling (RMS)
 - Utilization bound
- Optimal scheduling algorithm: Earliest Deadline First (EDF)
 - Utilization bound, implementation, evaluation
- Relaxing assumptions: when relative deadline $<$ period
 - Earliest Deadline First (EDF)
 - Processor demand analysis

Lab 4

- Lab 4 is due on 3/29
 - 6% in final grade
 - You will work on it on your own
 - No collaboration is allowed!
 - Need to let the TA check you off
- Real-Time Scheduling
 - <http://www.pitt.edu/~weigao/ece1175/spring2021/lab4.htm>