

ECE 1175  
Embedded System Design  
Real-Time Scheduling - II

Wei Gao

# Comparison

- **Schedulability**
  - RMS may not guarantee schedulability even when CPU is not fully utilized
  - EDF can guarantee schedulability as long as CPU is not fully utilized
- **Overhead**
  - RMS: low overhead, priorities are never changed
  - EDF: higher overhead, task priorities may need to be changed online
- **Optimality**
  - RMS: optimal **static** priority scheduling algorithm
  - EDF: optimal **dynamic** priority scheduling algorithm

# Relaxing Assumptions

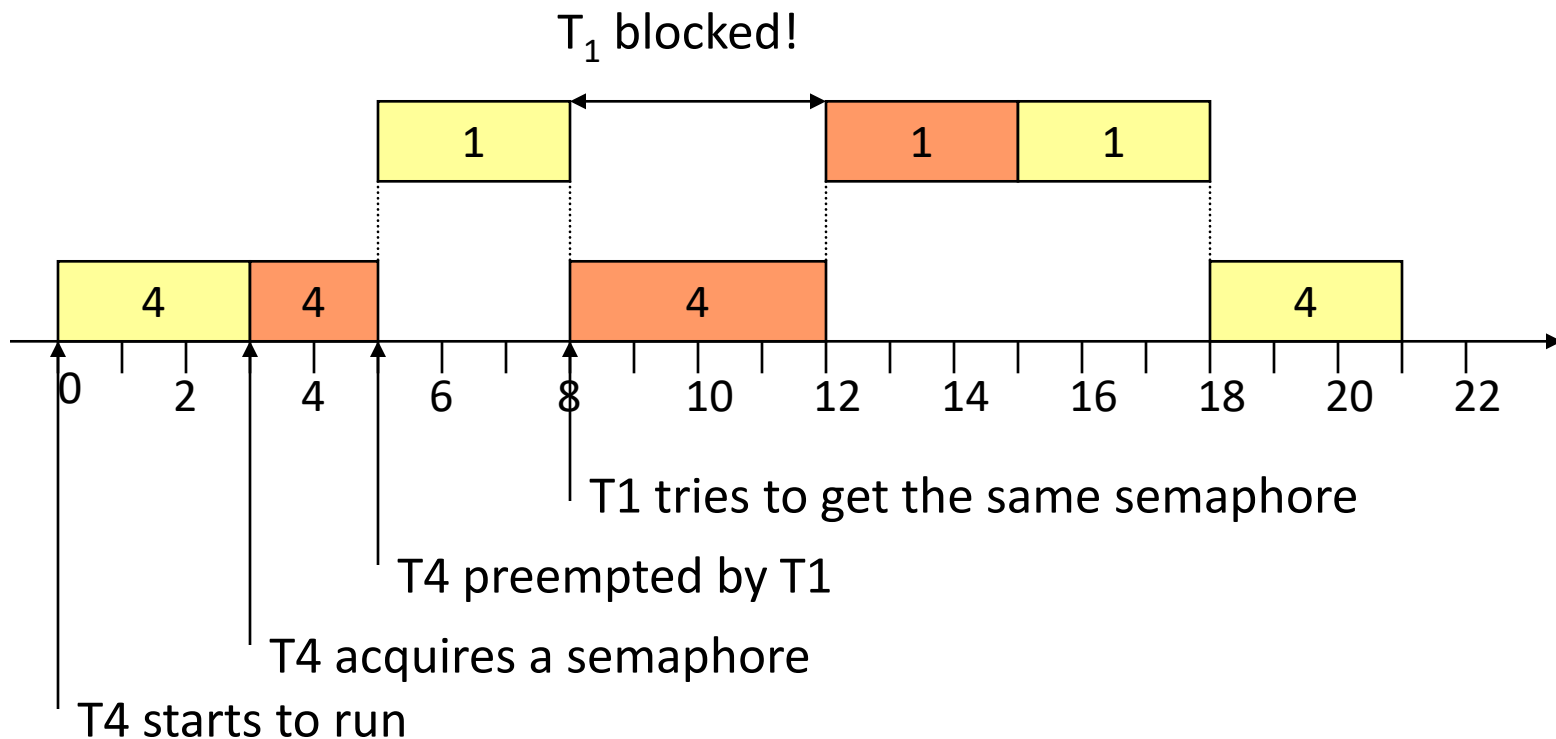
- Single processor.
  - All tasks are periodic.
  - Zero context switch time.
  - Relative deadline = period.
  - No priority inversion.
- 
- What if priority inversion exists?

# Priority Inversion


- A lower-priority task blocks a higher-priority task from running.
- Sources of priority inversion
  - Access shared resources guarded by semaphores
    - Lower-priority task gets the resource first
  - Access non-preemptive subsystems
    - Communication subsystems
    - Storage

# Priority Inversion

critical section




# Unbounded Priority Inversion



# Solution: Priority Inheritance

- Let the low-priority task inherit the priority of the blocked high-priority task.



# Real Incident

- Mars Pathfinder
- Priority inversion on Mars
  - [http://www.cse.chalmers.se/~risat/Report\\_MarsPathFinder.pdf](http://www.cse.chalmers.se/~risat/Report_MarsPathFinder.pdf)
  - <https://www.youtube.com/watch?v=lyx7kARrGeM>






# Relaxing Assumptions

- Single processor.
  - All tasks are periodic.
  - Zero context switch time.
  - Relative deadline = period.
  - No priority inversion. (relaxed)
- 
- What if we have multiple processors?


# End-to-End Task Model

- An (end-to-end) task is composed of multiple subtasks running on multiple processors
  - Message/event
  - Remote method invocation
- Subtasks are subject to precedence constraints
  - Task = a **chain**/tree/graph of subtasks
  - E.g. ship navigation



# Notation

- $T_i = \{T_{i,1}, T_{i,2}, \dots, T_{i,n(i)}\}$ 
  - $n(i)$ : the number of subtasks of  $T_i$
- Precedence constraint: Job  $J_{i,j}$  cannot be released until  $J_{i,j-1}$  is completed.



# End-to-End Scheduling Framework



1. **Task allocation**
2. Synchronization protocol
3. Subdeadline assignment
4. Schedulability analysis

# Task Allocation

- Load code (e.g., objects) to processors
- Strategies
  - Offline, static allocation subject to resource availability
  - Allocate a task when it arrives dynamically
  - Re-allocate (migrate) a task after it starts
- Optimal solutions for maximum schedulability
  - How to meet all deadlines in an optimal way?
    - E.g., minimize the number of needed processors
  - NP-hard: heuristics needed


# Bin Packing for Task Allocation

- Pack subtasks to bins (processors) with limited capacity
  - Size of a subtask
    - Utilization:  $C_{i,j}/P_i$
  - Capacity of each bin is its utilization bound
  - **Goal:** minimize the number of bins subject to the capacity constraints



# End-to-End Scheduling Framework


1. Task allocation
2. **Synchronization protocol**
3. Subdeadline assignment
4. Schedulability analysis



After a subtask is finished, should the next subtask start immediately or wait for a while?

# Greedy Protocol

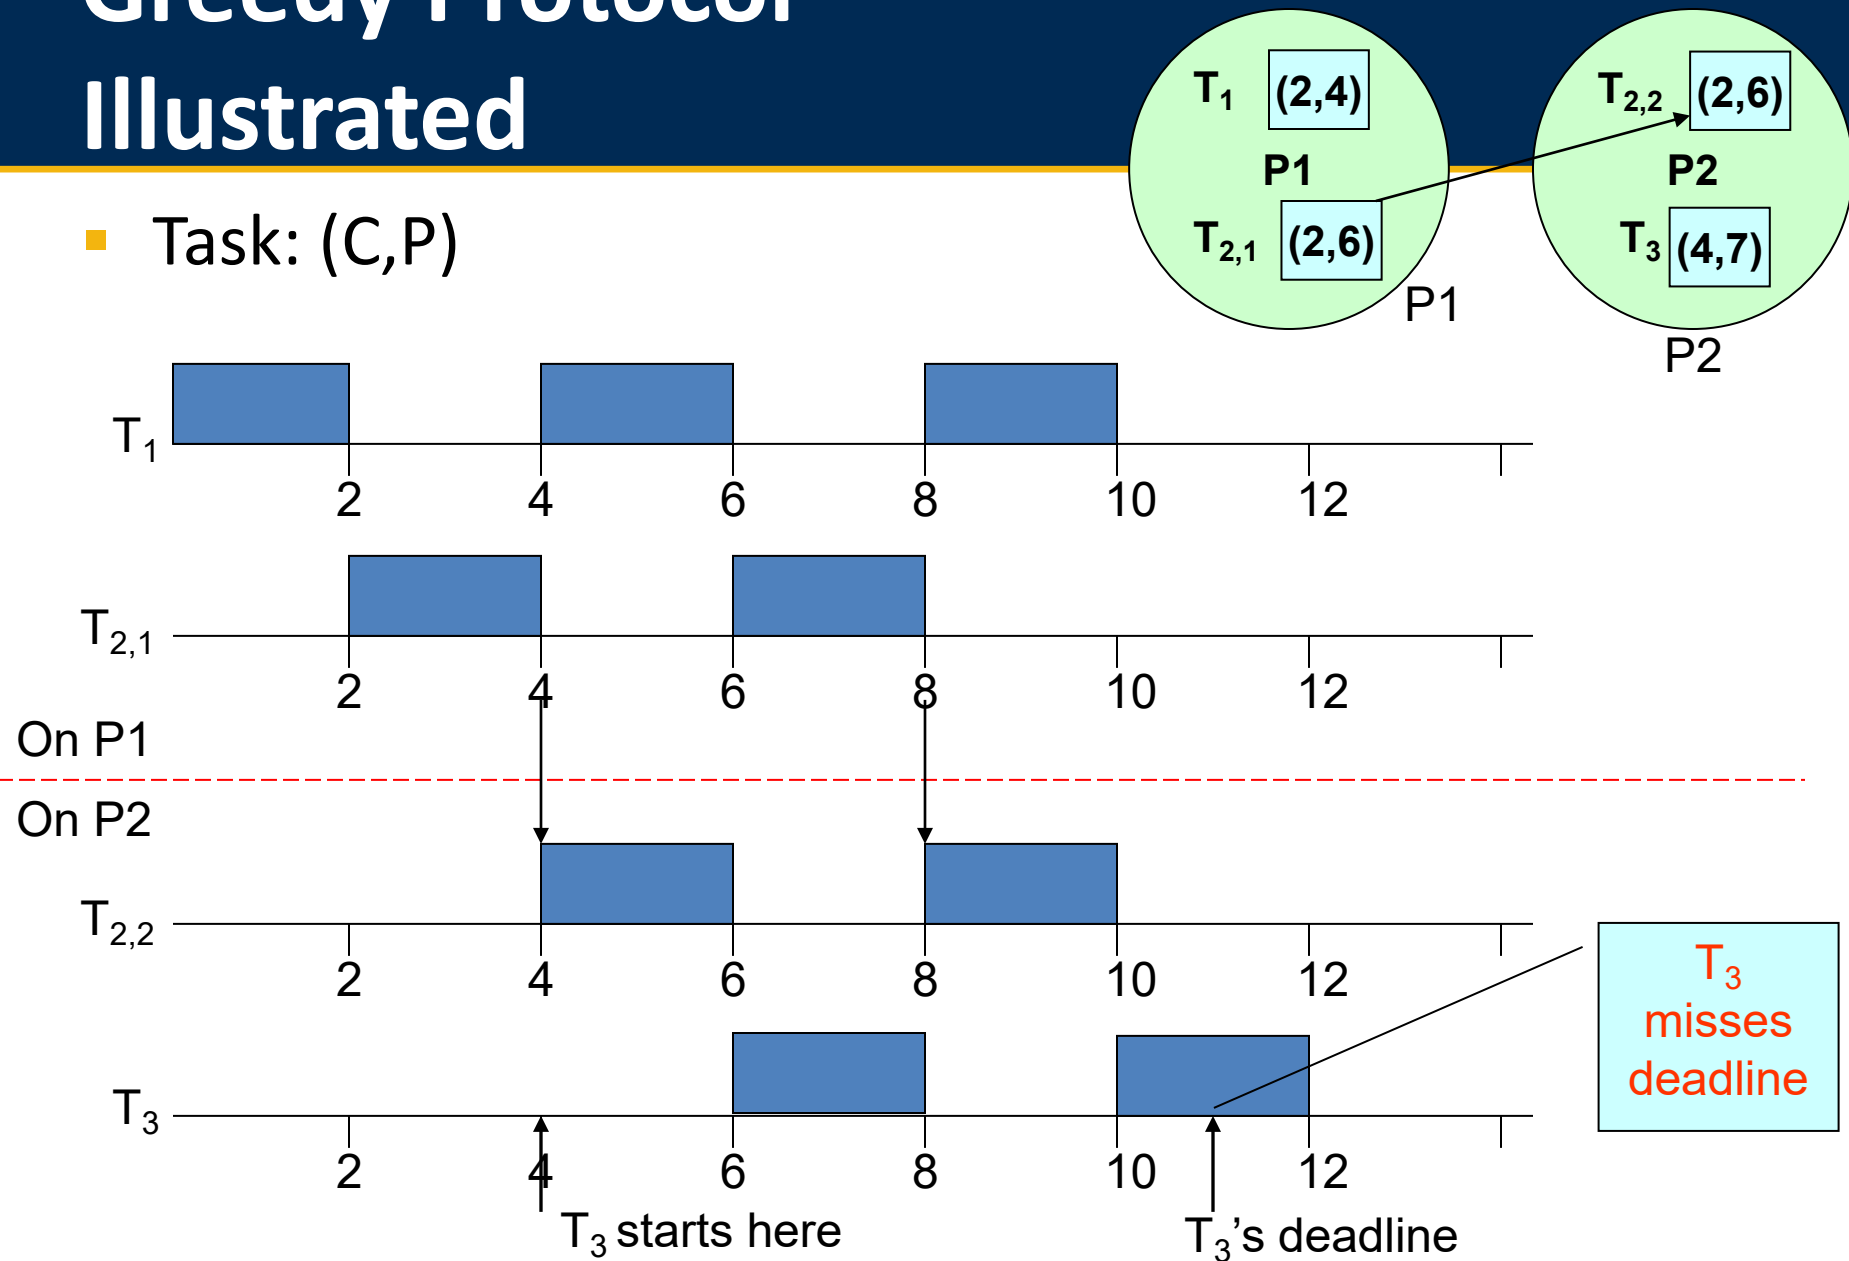
- After a subtask is finished, the next subtask starts **immediately**
- Release job  $J_{i,j;k}$  **as soon as**  $J_{i,j-1;k}$  is completed
- Subsequent subtasks may **not** be periodic under a greedy protocol
  - Difficult for schedulability analysis
  - High-priority tasks arrive early  $\rightarrow$  high worst-case response time for lower-priority tasks





# Greedy Protocol Illustrated

- Task: (C,P)



# Properties of Greedy Protocol

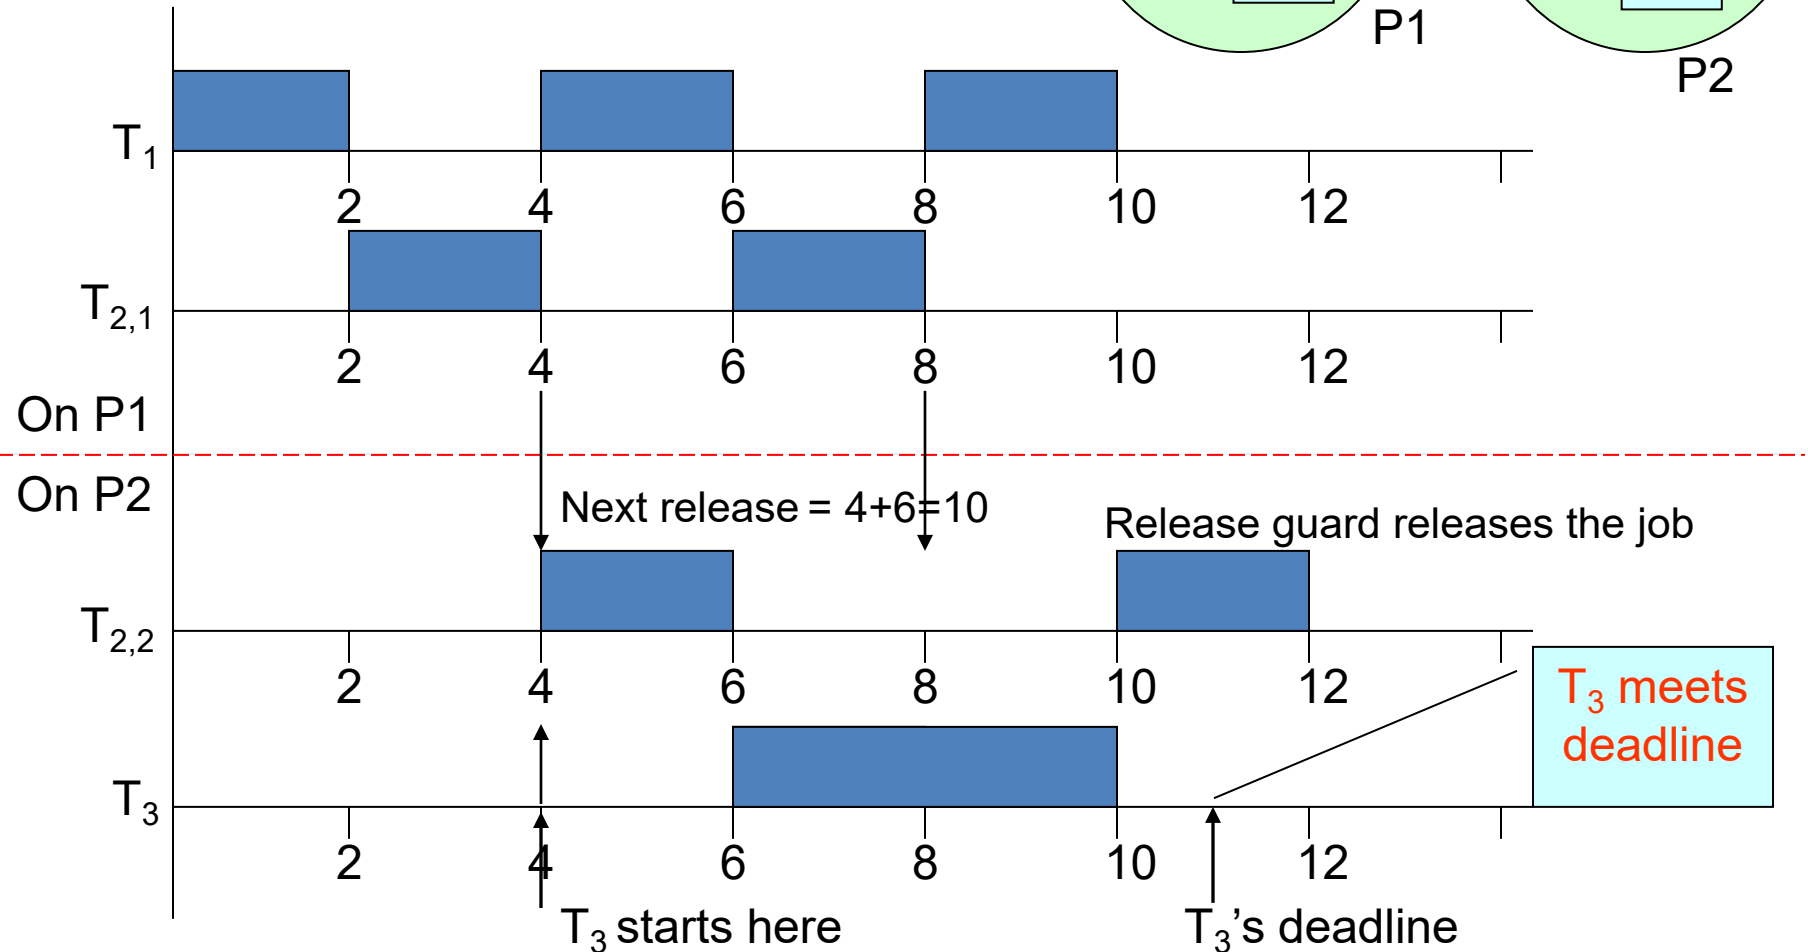

- Low overhead
- Low average response time
- Difficult schedulability analysis
  - Subsequent subtasks are no longer periodic
- High worst-case response time

# Release Guard

- After a subtask is finished, the next subtask may wait for a while before release
- Every subtask (if not a first subtask) has a release guard, which
  - waits for the preceding subtask for a result/event
  - then releases the job
    - at the point of exact one period from the last release time (Rule1)  
OR
    - whenever the processor becomes idle (Rule 2)
- Release guard strategy improves worst response time without affecting schedulability

# Release Guard Illustrated

## ■ Task: (C,P)



# End-to-End Scheduling Framework

1. Task allocation
2. Synchronization protocol
3. **Subdeadline assignment**
4. Schedulability analysis

# Subdeadline Assignment Algorithms

- Notation

- (Relative) deadline  $D_i$  of task  $T_i$
- (Relative) subdeadline  $D_{ij}$  of subtask  $T_{ij}$  ( $1 \leq j \leq n(i)$ )

- Ultimate Deadline (UD):  $D_{ij} = D_i$

- Example

- An end-to-end task  $T_1$ , with deadline as 12, has 4 subtasks:  $T_{11}$ ,  $T_{12}$ ,  $T_{13}$  and  $T_{14}$  with execution times as: 3, 1, 1, 1.
- $D_{11} = D_{12} = D_{13} = D_{14} = D_1 = 12$
- But  $T_{11}$ ,  $T_{12}$ ,  $T_{13}$  must finish earlier than the end-to-end deadline such that  $T_{14}$  can have time to run.

# Common Assignment Algorithms

- Proportional Deadline (PD):

- Assign deadline proportionally to execution time

$$D_{ij} = D_i \frac{C_{ij}}{\sum_{k=1}^{n(i)} C_{ik}}$$

- Example

- An end-to-end Task T1, with deadline as 12, has 4 subtasks T11, T12, T13 and T14 with execution time as: 3, 1, 1, 1.
- $D_{11} = 12 * 3/(3+1+1+1) = 6$
- $D_{12} = 12 * 1/(3+1+1+1) = 2$
- $D_{13} = 12 * 1/(3+1+1+1) = 2$
- $D_{14} = 12 * 1/(3+1+1+1) = 2$

# End-to-End Scheduling Framework

1. Task allocation
2. Synchronization protocol
3. Subdeadline assignment
4. **Schedulability analysis**
  - Decide an appropriate scheduling algorithm
    - RMS, EDF, etc
  - For each processor, conduct uniprocessor schedulability analysis
    - Then synchronize among multiple processors



# Summary

- End-to-end scheduling framework
  - Task allocation
    - Bin packing
  - Synchronization protocol
    - Greedy protocol, release guard
  - Subdeadline assignment
    - Ultimate deadline, proportional deadline
  - Schedulability analysis

# Reading

- Priority inversion on Mars
  - [http://www.cse.chalmers.se/~risat/Report\\_MarsPathFinder.pdf](http://www.cse.chalmers.se/~risat/Report_MarsPathFinder.pdf)
  - <https://www.youtube.com/watch?v=lyx7kARrGeM>